

CS 561: Data Systems Architectures

Tarikul Islam Papon, Zichen Zhu

papon@bu.edu, zczhu@bu.edu

<https://bu-disc.github.io/CS561/>

Tarikul Islam Papon



**PhD candidate @ Boston University
Lecturer @ BUET (Bangladesh University
of Engineering & Technology)**

<https://cs-people.bu.edu/papon/>
Office: CCDS 925 (TBD)
Office hours: Tu 9-10am

Zichen Zhu



**PhD candidate @ Boston University
MPhil @ HKU (the University of Hong
Kong)**

<https://cs-people.bu.edu/zczhu/>
Office: CCDS 925 (TBD)
Office hours: Th 1:30-2:30pm

Today

big data

data-driven world

data systems

which are the driving trends?

why do we need new designs?

CS 561 goals & logistics



I want you to speak up!
[and you can always interrupt me]

CS 561 philosophy

cutting-edge research

question everything (to understand it better!)

There are no *stupid questions!*

interactive & collaborative
projects, presentations, labs, OH




Understanding a design/system/algorithm ...

system

- component 1
- component 2
- component 3

algorithm

- step 1
- step 2
- step 3



understanding all steps and all decisions
helps us see the ***big picture***
and do **good research!**

(otherwise, we make ad hoc choices!)

Ask Questions!



... and answer my questions!


our **main goal** is to have **interesting discussions** that will help to gradually understand what the material discusses

(it's ok if not everything is clear, as long as you have questions!)

What do we do in this class?



reading papers



reviews



projects



Reading Papers

every class **1-2 papers to discuss** in detail

*in some classes the discussion will be led by a group of students
so that, each student will present one paper during the semester*

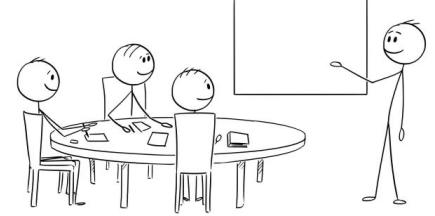
(background papers also available to provide more details)

read all of them!

write 4 reviews

answer one technical question per week (for a subset of the papers)

Presenting Papers



2-4 students will be responsible for presenting the paper
(discussing all main points of a review – see next slide)

during the presentation **anyone can ask questions** (including me!)
and each question is **addressed to all** (including me!)

prepare **slides** at least a **week before your presentation**



Writing Reviews

4 reviews and the **6 technical questions**
(some weeks will be “free”!)

learn

critic

review (up to one page)

what is the problem & why it is important?

why is it hard & why older approaches are not enough?

what is the key idea and why it works?

what is missing and how can we improve this idea?

does the paper support its claims?

possible next steps of the work presented in the paper?

single technical question

to make sure the heart of the paper is clearly understood

remember, this will help us do **good research!**

Projects

AND

project 0

A small implementation project
to sharpen dev skills

independent project



Due on Feb 2, 2024

project 1

A medium project to give you a flavor of
large-scale production system

groups of 3



Due on Feb 16, 2024


Projects

AND

project 0

A small implementation project
to sharpen dev skills

independent project



Due on Feb 2, 2024

project 1

A medium project to give you a flavor of
large-scale production system

groups of 3



Due on Feb 16, 2024

Projects

systems project

groups of 2/3

implementation-heavy C/C++ project

OR

research project

groups of 3

pick a subject (list will be available)

design & analysis

experimentation



Projects

systems project

groups of 2/3

implementation-heavy C/C++ project

OR

research project

groups of 3

pick a subject (list will be available)

design & analysis

experimentation



Research Project: open questions

skew-aware join optimization

context-aware spatial indexes

exploit *near-sorted data* with concurrency control

quantify *Write Amplification* in modern SSDs

come up with your **own topic!**

more on the website (soon)



A good project

- (1) has a clear plan by project proposal by **end-February** (5%)
- (2) has significant preliminary work done by **end-March** (5%)

evaluation at the **end of the semester** (30%)

- (i) present the key ideas of the implementation/new approach
- (ii) present a set of experiments supporting your claims

come to OH!

(more details for the projects in Class 4)



Class Goal

understand the internals of
data systems for data science

tune data systems through **adaptation** and **automation**

get acquainted with research in the area

Can I take this class?



background

C++ programming
data structures
algorithms
comp. architecture

pre-req

CS460/660 & CS210
contact Papon/Zichen if not sure

how to be sure?

if familiar with most, then maybe!
if familiar with **none**, then no!

Next classes

Class 1-2

logistics, big data, data systems, trends and outlook

Class 3

more basics on data systems, systems classification, graph, cloud

Class 4

intro to class project

Class 5 and beyond


present and **discuss** research papers from Papon/Zichen + students + guest lectures



big data?

who doesn't have a lot of data?

So what do we do with this data?




is data
analysis new?



? what is
really new?






Every day, we create 2.5 exabytes*
of data — 90% of the data in the
world today has been created in
the last two years alone.

[Understanding Big Data, IBM]

*exabyte = 10^9 GB



data management skills needed



100s of entries

pen & paper

$10^3\text{-}10^6$ of entries

UNIX tools and excel

10^9 of entries

custom solutions, programming

10^{12+} of entries

data systems

size (volume)

rate (velocity)

sources (variety)




big data

(it's not only about size)

all of the above plus ...

our ability to collect *machine-generated* data

scientific experiments




sensors

social



monitoring



micro-payments

Internet-of-Things



cloud



data analysis

*know what we
are looking for*



data exploration

*not sure what we
are looking for*






big data

data systems are
in the middle of this!

big data


**data
systems**



what is a data system?

a **data system** is a large software system
(a collection of algorithms and data structures)
that **stores data**, and provides the **interface** to
update and **access** them **efficiently**

the end goal is to make **data analysis** easy



*“relational databases
are the foundation of
western civilization”*

Bruce Lindsay, IBM Research
ACM SIGMOD Edgar F. Codd Innovations award 2012

**Big
Tech Era**

2019



+ growing need for tailored systems


Why?




more data



new hardware



new applications



new performance
goals

ORACLE®



Google

The big success of 5 decades of research

a declarative interface!

“ask and thou shall receive”




is this good?

ask **what** you want

data system


system decides **how**
to store & access



*“three things are important
in the database world:
**performance, performance,
and performance”***

Bruce Lindsay, IBM Research
ACM SIGMOD Edgar F. Codd Innovations award 2012

CS561: data systems kernel under the looking glass



this is where we will spend our time!

system architecture (row/column/hybrid)
indexing
relational/graph/key-value
scale-up/scale-out

goal: learn to design and implement a DB kernel

how to design a data system kernel?

what are its basic components?

algorithms/data structures/caching policies

what decisions should we make?


how to combine? how to optimize for hardware?

designing a DB kernel is complex

data system design complexity



application



performance



budget

thousands of options
millions of decisions
billions of combinations

let's think together: a simple DB kernel

a key-value system, each entry is a {key,value} pair

main operations: *put, get, scan, range scan, count*

workload has both reads (get, scan, range scan) *and writes (put)*

data

how to store and how to access data? 

how to efficiently delete?



designing a simple key-value system

what is the key/value?

are they stored together?

can read/write ratio change over time?

what to use? b-tree, hash-table, scans, skip-lists, zonemaps?

how to handle concurrent queries? million concurrent queries?

what happens if data does not fit in memory?

how to compress data?

what about privacy and security?

how to offer robustness guarantees?

what happens when we move to the cloud?



hardware at massive scale

performance tradeoffs different

10GB app: 1% less memory in your machine

so what?

10GB app: 1% less memory in 1M instances

$1M * 10GB * 1\% = 100TB!$

~800k\$ in today's price

class key goal

understand system design tradeoffs

design and prototype a system

with other **side-effects**:

sharpening your systems skills

(C/C++, profiling, debugging, linux tools)

data system designer & researcher
any business, any startup, any scientific domain

CS 561: more logistics

topics

storage layouts, HTAP systems, adaptive indexing, solid-state storage, data integration, data skipping, data systems and ML, learned index

past but still relevant topics

relational systems, row-stores, query optimization, concurrency control, SQL

no textbook – only research papers

grading



class participation: 5%
project 0: 10%
project 1: 15%
reviews: 7%
technical questions: 8%
paper presentation: 15%
project proposal: 5%
mid-semester project report: 5%
project: 30%

Piazza



all discussions & announcements

<https://piazza.com/bu/spring2024/cs561>

also available on class website

How can I prepare?

1) Read background research material

- **Architecture of a Database System.** By J. Hellerstein, M. Stonebraker and J. Hamilton. Foundations and Trends in Databases, 2007
- **The Design and Implementation of Modern Column-store Database Systems.** By D. Abadi, P. Boncz, S. Harizopoulos, S. Idreos, S. Madden. Foundations and Trends in Databases, 2013

2) Start going over the papers

class summary

2 classes + 2 OH + 1 Lab (5 days) per week

each student

1 review/question per week + 1 paper presentation

project 0 + project 1 + *systems or research project*



proposal + mid-semester report + final report + project presentation

what to do now?

- A) **read the syllabus and the website**
- B) **register to Piazza + Gradescope**
- C) **start working on project 0**
- D) register for the presentation (week 2)
- E) start submitting paper reviews/answering tech. questions (week 3)
- F) go over the class project (end of next week will be available)
- G) start working on the proposal (week 3)

survival guide

class website: <https://bu-disc.github.io/CS561/>

piazza website: <https://piazza.com/bu/spring2024/cs561>

presentation registration: <http://tinyurl.com/S24-CS561-presentations>

gradescope: <https://www.gradescope.com/courses/693490> (**code in Piazza**)

office hours: Papon (Tu 9-10am)

Zichen (Th 1:30-2:30pm)

material: papers available from the BU network

Welcome to CS 561: Data Systems Architectures!

Tarikul Islam Papon, Zichen Zhu

papon@bu.edu, zczhu@bu.edu

next time: more detailed logistics and start with data systems design